

Heart Disease Classifier

September 30, 2024

1 Heart Disease

Project By Mohamed Ouhami

In this document, we will apply the different concepts and technologies learned to work on classifying heart diseases

The data is coming from the UC Irvine website, and contains approximately 14 features and 1 target

1.1 Framing the problem

1.1.1 Defining the objective in business terms

Our main objective is to create a Machine learning model that can classify whether a patient has a heart disease based on many features containing info about his life style going from his heart rate to his smoking habits

Only 14 attributes used: 1. #3 (age)

2. #4 (sex)

3. #9 (cp) cp: chest pain type – Value 1: typical angina – Value 2: atypical angina – Value 3: non-anginal pain – Value 4: asymptomatic

4. #10 (trestbps) resting blood pressure (in mm Hg on admission to the hospital)

5. #12 (chol)

6. #16 (fbs) (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)

7. #19 (restecg) resting electrocardiographic results – Value 0: normal – Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV) – Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria

8. #32 (thalach) maximum heart rate achieved

9. #38 (exang) exercise induced angina (1 = yes; 0 = no)

10. #40 (oldpeak) ST depression induced by exercise relative to rest

11. #41 (slope) the slope of the peak exercise ST segment – Value 1: upsloping – Value 2: flat – Value 3: downsloping

12. #44 (ca) number of major vessels (0-3) colored by flourosopy

13. #51 (thal) thal: 3 = normal; 6 = fixed defect; 7 = reversable defect 14. #58 (num) (the predicted attribute)

1.1.2 How will our solution be used ?

Our solution will be used in a desktop or web application in hospitals that the doctors are gonna use to help them see the probability of the patient having a heart disease

We're going to create an API that can be used to introduce the patient info into It, to be able to make predictions on whether the patient has a disease or not

The API will be made using FastAPI to create endpoints ##

1.2 How we would frame the problem ?

This problem can be framed as being a classification supervised learning offline problem, as our goal is to be able to classify the patients on a pre-established database that has labels on the data

1.3 How should the performance be measured ?

As this is a classification problem, we will go for creating a confusion matrix, and calculate the precision and recall on them, to ensure that model gives reasonable answers, and to see how can we improve on some instances of the target

1.4 Is the performance measure aligned with the business objective?

The performance measure is indeed aligned with the business objective, as we can get getting false positives, than having false negatives, as It can lead to the patient's death

1.5 What would be the minimum performance needed to reach the business objective

The minimum performance would be to have as much precision for true positives, and as lowest rate of false negatives, We would opt for 95%

1.6 Is human expertise available

Yes, the measurements should be done by a professional to get as closest precision as possible and for the measurements to not be noisy or mis-conducted

1.7 Get the data

We can automatically download the data from the uci using their API using the code below with the id of the repo

```
[2]: import pandas as pd
import numpy as np
```

```
[3]: from ucimlrepo import fetch_ucirepo

# fetch dataset
heart_disease = fetch_ucirepo(id=45)
```

Now, we can go and see some information about the data including the first 5 rows of the data, description and info about It

```
[4]: heart_disease.data
```

```
[4]: {'ids': None,
      'features':      age sex cp trestbps chol fbs restecg thalach exang
oldpeak \
0      63      1      1      145      233      1      2      150      0      2.3
1      67      1      4      160      286      0      2      108      1      1.5
2      67      1      4      120      229      0      2      129      1      2.6
3      37      1      3      130      250      0      0      187      0      3.5
4      41      0      2      130      204      0      2      172      0      1.4
..      ...      ...      ..      ...      ...      ...      ...      ...
298    45      1      1      110      264      0      0      132      0      1.2
299    68      1      4      144      193      1      0      141      0      3.4
300    57      1      4      130      131      0      0      115      1      1.2
301    57      0      2      130      236      0      2      174      0      0.0
302    38      1      3      138      175      0      0      173      0      0.0
```

```
      slope      ca      thal
0          3      0.0      6.0
1          2      3.0      3.0
2          2      2.0      7.0
3          3      0.0      3.0
4          1      0.0      3.0
..      ...      ...      ...
298      2      0.0      7.0
299      2      2.0      7.0
300      2      1.0      7.0
301      2      1.0      3.0
302      1      NaN      3.0
```

```
[303 rows x 13 columns],
```

```
'targets':      num
0          0
1          2
2          1
3          0
4          0
..      ...
298      1
299      2
300      3
301      1
302      0
```

```
[303 rows x 1 columns],
```

```
'original':      age sex cp trestbps chol fbs restecg thalach exang
oldpeak \
0      63      1      1      145      233      1      2      150      0      2.3
1      67      1      4      160      286      0      2      108      1      1.5
```

2	67	1	4	120	229	0	2	129	1	2.6
3	37	1	3	130	250	0	0	187	0	3.5
4	41	0	2	130	204	0	2	172	0	1.4
..
298	45	1	1	110	264	0	0	132	0	1.2
299	68	1	4	144	193	1	0	141	0	3.4
300	57	1	4	130	131	0	0	115	1	1.2
301	57	0	2	130	236	0	2	174	0	0.0
302	38	1	3	138	175	0	0	173	0	0.0

	slope	ca	thal	num
0	3	0.0	6.0	0
1	2	3.0	3.0	2
2	2	2.0	7.0	1
3	3	0.0	3.0	0
4	1	0.0	3.0	0
..
298	2	0.0	7.0	1
299	2	2.0	7.0	2
300	2	1.0	7.0	3
301	2	1.0	3.0	1
302	1	NaN	3.0	0

```
[303 rows x 14 columns],
'headers': Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg',
'thalach',
'exang', 'oldpeak', 'slope', 'ca', 'thal', 'num'],
dtype='object')}]
```

```
[5]: # Copy the data into a df
df = heart_disease.data.original.copy()
columns_names = heart_disease.data.headers.tolist()

column_target = columns_names.pop(-1)
columns_features = columns_names.copy()
```

```
[6]: print(f"These are the feature columns {columns_features}")
print(f"This is the target column {column_target}")
```

These are the feature columns ['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal']
This is the target column num

Now, we will go and study each attribute of the data

```
[7]: df.head()
```

```
[7]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	1	145	233	1	2	150	0	2.3	3	
1	67	1	4	160	286	0	2	108	1	1.5	2	
2	67	1	4	120	229	0	2	129	1	2.6	2	
3	37	1	3	130	250	0	0	187	0	3.5	3	
4	41	0	2	130	204	0	2	172	0	1.4	1	

	ca	thal	num
0	0.0	6.0	0
1	3.0	3.0	2
2	2.0	7.0	1
3	0.0	3.0	0
4	0.0	3.0	0

```
[8]: # Getting numerical statistics of our data
df.describe()
```

```
[8]:
```

	age	sex	cp	trestbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.438944	0.679868	3.158416	131.689769	246.693069	0.148515	
std	9.038662	0.467299	0.960126	17.599748	51.776918	0.356198	
min	29.000000	0.000000	1.000000	94.000000	126.000000	0.000000	
25%	48.000000	0.000000	3.000000	120.000000	211.000000	0.000000	
50%	56.000000	1.000000	3.000000	130.000000	241.000000	0.000000	
75%	61.000000	1.000000	4.000000	140.000000	275.000000	0.000000	
max	77.000000	1.000000	4.000000	200.000000	564.000000	1.000000	

	restecg	thalach	exang	oldpeak	slope	ca	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	299.000000	
mean	0.990099	149.607261	0.326733	1.039604	1.600660	0.672241	
std	0.994971	22.875003	0.469794	1.161075	0.616226	0.937438	
min	0.000000	71.000000	0.000000	0.000000	1.000000	0.000000	
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000	
50%	1.000000	153.000000	0.000000	0.800000	2.000000	0.000000	
75%	2.000000	166.000000	1.000000	1.600000	2.000000	1.000000	
max	2.000000	202.000000	1.000000	6.200000	3.000000	3.000000	

	thal	num
count	301.000000	303.000000
mean	4.734219	0.937294
std	1.939706	1.228536
min	3.000000	0.000000
25%	3.000000	0.000000
50%	3.000000	0.000000
75%	7.000000	2.000000
max	7.000000	4.000000

```
[9]: # Getting info about the columns and their data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null   int64
 1   sex         303 non-null   int64
 2   cp          303 non-null   int64
 3   trestbps    303 non-null   int64
 4   chol        303 non-null   int64
 5   fbs         303 non-null   int64
 6   restecg     303 non-null   int64
 7   thalach     303 non-null   int64
 8   exang       303 non-null   int64
 9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          299 non-null   float64
12  thal        301 non-null   float64
13  num         303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

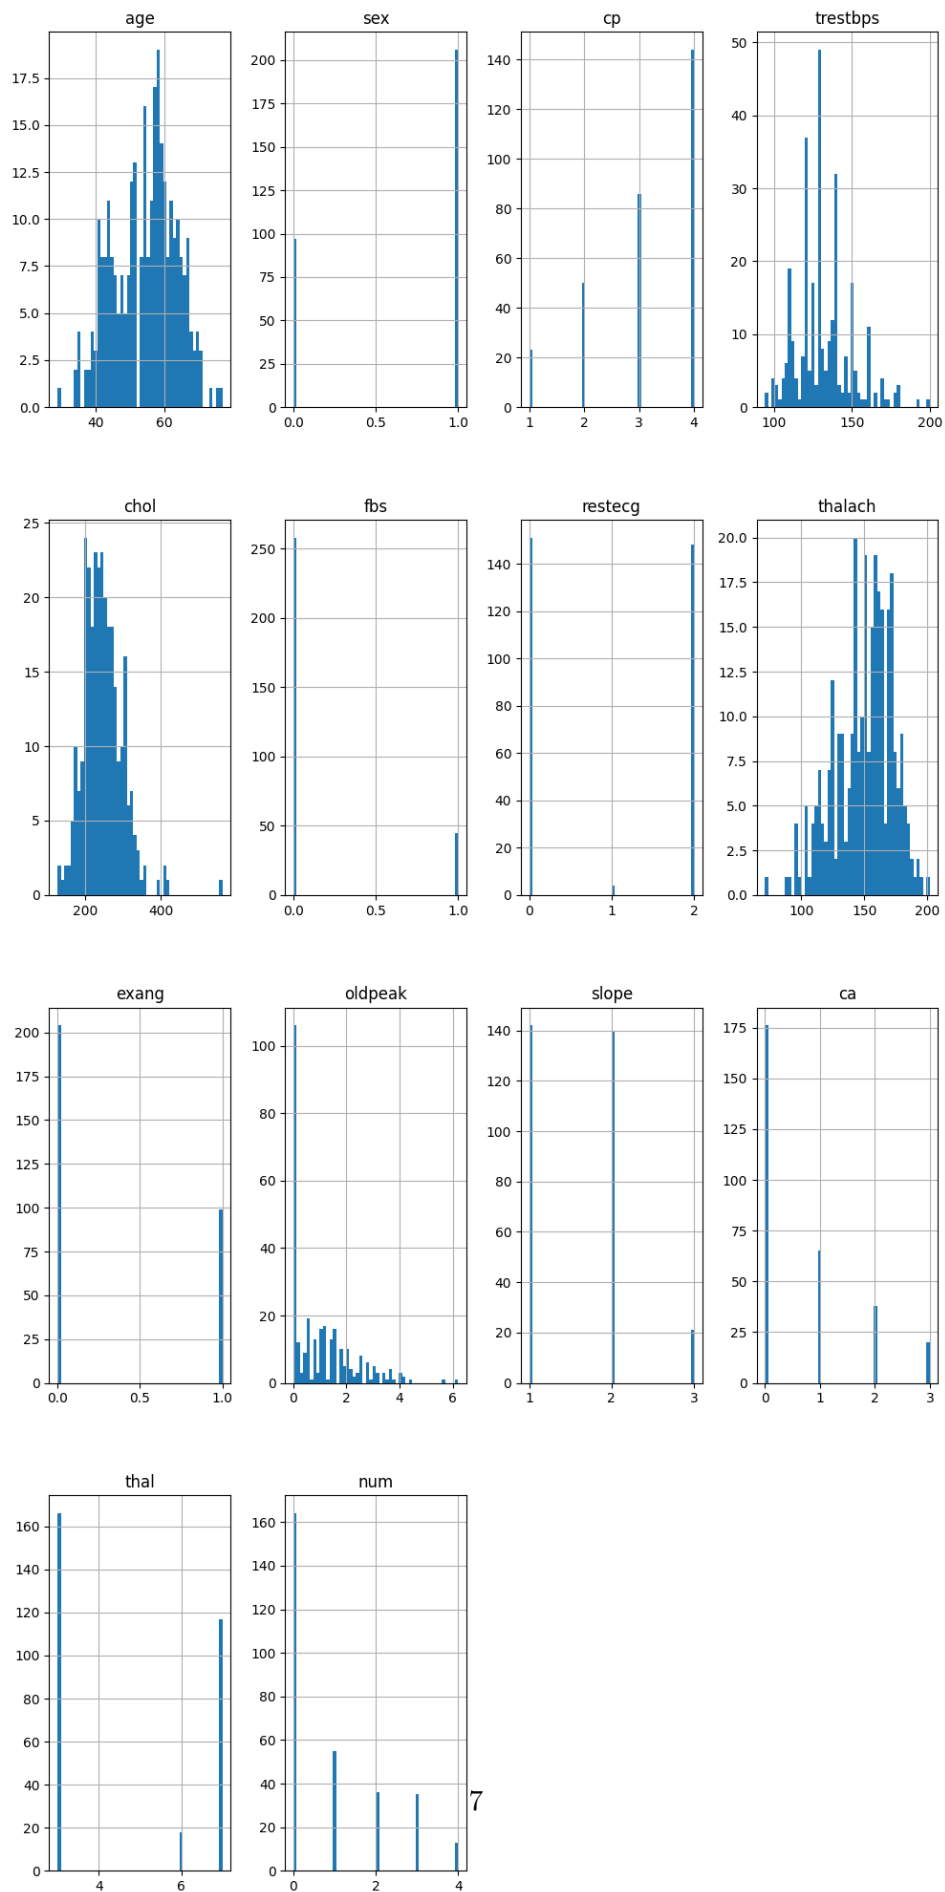
We can see that the data contain 303 entries with 14 columns, 13 of which are features and one target with some missing values in the ca and thal

Now, we will go and plot the hist plot to see the distribution of the data

```
[10]: import matplotlib.pyplot as plt
```

```
[11]: df.hist(bins=50,figsize=(12,25))
```

```
[11]: array([[<Axes: title={'center': 'age'}>, <Axes: title={'center': 'sex'}>,
  <Axes: title={'center': 'cp'}>,
  <Axes: title={'center': 'trestbps'}>],
  [<Axes: title={'center': 'chol'}>,
  <Axes: title={'center': 'fbs'}>,
  <Axes: title={'center': 'restecg'}>,
  <Axes: title={'center': 'thalach'}>],
  [<Axes: title={'center': 'exang'}>,
  <Axes: title={'center': 'oldpeak'}>,
  <Axes: title={'center': 'slope'}>,
  <Axes: title={'center': 'ca'}>],
  [<Axes: title={'center': 'thal'}>,
  <Axes: title={'center': 'num'}>, <Axes: >, <Axes: >]],
  dtype=object)
```



We can see that most of the features has a normal distribution with a bell-like shape around the mean for each numerical data, except for the categorical data

For the target column called “num”, It represents diagnosis of heart disease (angiographic disease status)

- Value 0: < 50% diameter narrowing
- Value 1: > 50% diameter narrowing

```
[12]: df[column_target].astype("category").value_counts
```

```
[12]: <bound method IndexOpsMixin.value_counts of 0      0
      1      2
      2      1
      3      0
      4      0
      ..
     298      1
     299      2
     300      3
     301      1
     302      0
      Name: num, Length: 303, dtype: category
      Categories (5, int64): [0, 1, 2, 3, 4]>
```

1.8 Preparing the data

1.8.1 Cleaning the data

We will clean the data by removing values, outliers and all

```
[13]: df_copy = df.copy()
```

```
[14]: # We will go and fix the missing values, removing them is the best options as
      ↪ the missing values are in categorical features and not numerical
      df_copy = df_copy.dropna(axis=0)

      df_copy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 297 entries, 0 to 301
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         297 non-null    int64
1   sex         297 non-null    int64
2   cp          297 non-null    int64
3   trestbps    297 non-null    int64
```



```

4   chol      297 non-null   int64
5   fbs       297 non-null   int64
6   restecg   297 non-null   int64
7   thalach   297 non-null   int64
8   exang     297 non-null   int64
9   oldpeak   297 non-null   float64
10  slope     297 non-null   int64
11  ca        297 non-null   float64
12  thal      297 non-null   float64
13  num       297 non-null   int64

```

dtypes: float64(3), int64(11)

memory usage: 34.8 KB

```

[15]: # Feature scaling
from sklearn.preprocessing import MinMaxScaler

numerical_columns = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']

df_copy_numerical = df_copy[numerical_columns]

min_max_scaler = MinMaxScaler()
data_scaled = min_max_scaler.fit_transform(df_copy_numerical)

```

```

[16]: df_scaled = pd.DataFrame(data_scaled, columns=numerical_columns)
df_scaled

```

```

[16]:
      age  trestbps    chol  thalach  oldpeak
0   0.708333  0.481132  0.244292  0.603053  0.370968
1   0.791667  0.622642  0.365297  0.282443  0.241935
2   0.791667  0.245283  0.235160  0.442748  0.419355
3   0.166667  0.339623  0.283105  0.885496  0.564516
4   0.250000  0.339623  0.178082  0.770992  0.225806
..      ...      ...      ...      ...      ...
292  0.583333  0.433962  0.262557  0.396947  0.032258
293  0.333333  0.150943  0.315068  0.465649  0.193548
294  0.812500  0.471698  0.152968  0.534351  0.548387
295  0.583333  0.339623  0.011416  0.335878  0.193548
296  0.583333  0.339623  0.251142  0.786260  0.000000

```

[297 rows x 5 columns]

Now, we're going to delete the old columns , and add the new one

```

[17]: df_copy = df_copy.drop(columns=numerical_columns)

```

```

[18]: df_copy_scaled = pd.concat([df_copy, df_scaled], axis=1)

```

```
[19]: df_copy_scaled = df_copy_scaled.dropna(axis=0)
```

```
[20]: df_copy_scaled.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 292 entries, 0 to 296
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   sex         292 non-null    float64
1   cp          292 non-null    float64
2   fbs         292 non-null    float64
3   restecg     292 non-null    float64
4   exang       292 non-null    float64
5   slope       292 non-null    float64
6   ca          292 non-null    float64
7   thal        292 non-null    float64
8   num         292 non-null    float64
9   age         292 non-null    float64
10  trestbps    292 non-null    float64
11  chol        292 non-null    float64
12  thalach     292 non-null    float64
13  oldpeak     292 non-null    float64
dtypes: float64(14)
memory usage: 34.2 KB
```

```
[21]: df_copy_scaled
```

```
[21]:
```

	sex	cp	fbs	restecg	exang	slope	ca	thal	num	age	trestbps	\
0	1.0	1.0	1.0	2.0	0.0	3.0	0.0	6.0	0.0	0.708333	0.481132	
1	1.0	4.0	0.0	2.0	1.0	2.0	3.0	3.0	2.0	0.791667	0.622642	
2	1.0	4.0	0.0	2.0	1.0	2.0	2.0	7.0	1.0	0.791667	0.245283	
3	1.0	3.0	0.0	0.0	0.0	3.0	0.0	3.0	0.0	0.166667	0.339623	
4	0.0	2.0	0.0	2.0	0.0	1.0	0.0	3.0	0.0	0.250000	0.339623	
..	
292	1.0	4.0	0.0	0.0	1.0	3.0	0.0	6.0	2.0	0.583333	0.433962	
293	1.0	4.0	0.0	2.0	1.0	1.0	2.0	7.0	2.0	0.333333	0.150943	
294	0.0	4.0	0.0	0.0	1.0	2.0	0.0	3.0	1.0	0.812500	0.471698	
295	1.0	2.0	0.0	0.0	0.0	1.0	0.0	3.0	0.0	0.583333	0.339623	
296	1.0	4.0	1.0	2.0	0.0	2.0	2.0	6.0	3.0	0.583333	0.339623	
	chol	thalach	oldpeak									
0	0.244292	0.603053	0.370968									
1	0.365297	0.282443	0.241935									
2	0.235160	0.442748	0.419355									
3	0.283105	0.885496	0.564516									
4	0.178082	0.770992	0.225806									
..									

```

292  0.262557  0.396947  0.032258
293  0.315068  0.465649  0.193548
294  0.152968  0.534351  0.548387
295  0.011416  0.335878  0.193548
296  0.251142  0.786260  0.000000

```

```
[292 rows x 14 columns]
```

1.9 Testing models and picking a dirty model

In this part, we will try and test some models to see the accuracy and performance of each one before going fine tuning the model

1.9.1 Creating training and test dataset

```

[22]: from sklearn.model_selection import train_test_split

X = df_copy_scaled[columns_features]
y = df_copy_scaled[column_target]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.
↪2, random_state=42)

```

1.10 Testing some random models

As we're dealing with a classification problem, we will go with one of the known models (Logistic Regression, KNN, SVM, Decision Tree, RandomForest, GBM, AdaBoost, XGBoost)

```

[23]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

lg_reg = LogisticRegression()
knn_clf = KNeighborsClassifier()
svm_clf = SVC()
dt_clf = DecisionTreeClassifier()
rf_clf = RandomForestClassifier()

models = [lg_reg, knn_clf, svm_clf, dt_clf, rf_clf]

```

1.10.1 Creating a stratified dataset for better training

```

[24]: from sklearn.model_selection import StratifiedKFold
from sklearn.base import clone

# Assuming models is a list of initialized models

```

```

skfolds = StratifiedKFold(n_splits=3) # Specify the number of folds

# Perform stratified k-fold cross-validation
for train_index, test_index in skfolds.split(X_train, y_train):

    X_train_folds = X_train.iloc[train_index]
    X_test_folds = X_train.iloc[test_index]
    y_train_folds = y_train.iloc[train_index]
    y_test_folds = y_train.iloc[test_index]

    for model in models:
        model = clone(model) # Clone the model to avoid fitting on previous
        ↪folds

        model.fit(X_train_folds, y_train_folds) # Train the model

        y_hat = model.predict(X_test_folds) # Predict on the test fold

        n_correct = sum(y_hat == y_test_folds.values) # Compare predictions
        ↪with the actual test set

        precision = n_correct / len(y_test_folds) # Calculate precision
        ↪(accuracy for this fold)

        print(f"For the model {model.__class__.__name__}, we got a precision of
        ↪{precision:.2f}")

```

c:\Users\moham\miniconda3\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

For the model LogisticRegression, we got a precision of 0.60

For the model KNeighborsClassifier, we got a precision of 0.60

For the model SVC, we got a precision of 0.62

For the model DecisionTreeClassifier, we got a precision of 0.51

For the model RandomForestClassifier, we got a precision of 0.65

For the model LogisticRegression, we got a precision of 0.63

For the model KNeighborsClassifier, we got a precision of 0.64

For the model SVC, we got a precision of 0.59

For the model DecisionTreeClassifier, we got a precision of 0.53

```
c:\Users\moham\miniconda3\Lib\site-  
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

For the model RandomForestClassifier, we got a precision of 0.63

For the model LogisticRegression, we got a precision of 0.60

For the model KNeighborsClassifier, we got a precision of 0.52

For the model SVC, we got a precision of 0.58

For the model DecisionTreeClassifier, we got a precision of 0.45

```
c:\Users\moham\miniconda3\Lib\site-  
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

For the model RandomForestClassifier, we got a precision of 0.52

Now, we're going to use Cross Val Evaluation to create the confusion matrix, and check the response of each model

```
[25]: import warnings  
  
# Suppress specific warnings  
warnings.filterwarnings("ignore", category=UserWarning, module='sklearn')
```

```
[27]: from sklearn.metrics import f1_score, precision_score, recall_score  
  
def calculate_scores(y_pred, y):  
  
    precisionScore = precision_score(y_pred, y, average='weighted')  
    recallScore = recall_score(y_pred, y, average='weighted')  
    f1Score = f1_score(y_pred, y, average='weighted')
```

```

print(f"The precision score is {precisionScore}\n")
print(f"The recall score is {recallScore}\n")
print(f"The f1score score is {f1Score}\n")

return precisionScore,recallScore,f1Score

```

```

[28]: from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix

model_scores = []
for model in models:

    y_train_pred = cross_val_predict(model,X_train,y_train,cv=3)
    print(f"For the model {model} : {confusion_matrix(y_train,y_train_pred)}")

    precisionScore,recallScore,f1Score = calculate_scores(y_train_pred,y_train)

    model_scores.append({'model':model.__class__.__name__,'score':recallScore})

```

```

For the model LogisticRegression() : [[115   8   1   3   0]
 [ 21  13   8   1   0]
 [  7   6  11   5   0]
 [  5   5  12   3   0]
 [  2   2   1   4   0]]

```

The precision score is 0.6890261064390364

The recall score is 0.6094420600858369

The f1score score is 0.6441218745776925

```

For the model KNeighborsClassifier() : [[115   7   3   2   0]
 [ 23  10   5   4   1]
 [  9  10   3   7   0]
 [  3   7   5   9   1]
 [  1   2   2   4   0]]

```

The precision score is 0.6709289504435193

The recall score is 0.5879828326180258

The f1score score is 0.6245335816877707

```

For the model SVC() : [[119   8   0   0   0]
 [ 23  20   0   0   0]
 [  8  21   0   0   0]
 [  8  17   0   0   0]
 [  2   7   0   0   0]]

```

The precision score is 0.7891620095047754

The recall score is 0.5965665236051502

The f1score score is 0.6774907322848573

```
For the model DecisionTreeClassifier() : [[89 23 10  3  2]
 [21  7  5  8  2]
 [ 9  6  9  4  1]
 [ 6  6  5  7  1]
 [ 1  3  2  2  1]]
```

The precision score is 0.4838766584837949

The recall score is 0.48497854077253216

The f1score score is 0.48428950691033174

```
For the model RandomForestClassifier() : [[119  7  0  1  0]
 [ 23  9  8  3  0]
 [ 11  9  5  4  0]
 [ 11  4  6  4  0]
 [  3  2  4  0  0]]
```

The precision score is 0.7246961557726715

The recall score is 0.5879828326180258

The f1score score is 0.6426982328269882

1.11 We can see that there is a sampling error, as the class 4 is not well represented, and is rarely predicted, so that would make our model biased towards the oversampled classes

A way to fix this is to do some data augmentation

Now, we're going to go and calculate the scores for the confusion matrix

```
[29]: print(max(model_scores, key=lambda x: x['score']))
```

```
{'model': 'LogisticRegression', 'score': np.float64(0.6094420600858369)}
```

Now, we can see the model with the highest accuracy is the Support Vector Machine with a precision of 67% without any fine-tuning parameters

Now, we'll use the grid search to get the best parameters combination for the Support Vector Machines

```
[30]: from sklearn.model_selection import GridSearchCV
      from sklearn.linear_model import LogisticRegression

      # Define the parameter grid for Logistic Regression
```

```

param_grid = [
    {
        'C': [0.01, 0.1, 1, 10, 100],
        'penalty': ['l2'], # 'l2' can be used with both 'liblinear' and 'lbfgs'
        'solver': ['liblinear', 'lbfgs'] # Both solvers support 'l2'
    },
    {
        'C': [0.01, 0.1, 1, 10, 100],
        'penalty': ['l1'], # 'l1' can only be used with 'liblinear' or 'saga'
        'solver': ['liblinear'] # Only 'liblinear' supports 'l1'
    },
    {
        'C': [0.01, 0.1, 1, 10, 100],
        'penalty': ['elasticnet'], # 'elasticnet' can only be used with 'saga'
        'solver': ['saga'], # Only 'saga' supports 'elasticnet'
        'l1_ratio': [0.1, 0.5, 0.9] # Include l1_ratio for elasticnet
    }
]

# Create an instance of the Logistic Regression model
logistic_clf_final = LogisticRegression(max_iter=1000) # Increase max_iter if
↳needed

# Set up GridSearchCV with recall_weighted as the scoring metric
grid_search = GridSearchCV(logistic_clf_final, param_grid, cv=3,
↳scoring='recall_weighted', error_score='raise')

# Fit the model to your training data
# grid_search.fit(X_train, y_train)

# You can now access the best parameters found by GridSearch
# best_params = grid_search.best_params_

```

```
[31]: grid_search.fit(X_train,y_train)
```

```

[31]: GridSearchCV(cv=3, error_score='raise',
        estimator=LogisticRegression(max_iter=1000),
        param_grid=[{'C': [0.01, 0.1, 1, 10, 100], 'penalty': ['l2'],
                        'solver': ['liblinear', 'lbfgs']},
                    {'C': [0.01, 0.1, 1, 10, 100], 'penalty': ['l1'],
                        'solver': ['liblinear']},
                    {'C': [0.01, 0.1, 1, 10, 100],
                        'l1_ratio': [0.1, 0.5, 0.9],
                        'penalty': ['elasticnet'], 'solver': ['saga']}],
        scoring='recall_weighted')

```

```
[32]: grid_search.best_params_
```



```
[32]: {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}
```

```
[33]: import joblib

final_model = grid_search.best_estimator_
joblib.dump(final_model, "heart_classifier.pkl")

y_hat = final_model.predict(X_test)

confusion_matrix(y_hat, y_test)

calculate_scores(y_hat, y_test)
```

The precision score is 0.7131655538435199

The recall score is 0.6101694915254238

The f1score score is 0.651964899120692

```
[33]: (np.float64(0.7131655538435199),
      np.float64(0.6101694915254238),
      np.float64(0.651964899120692))
```

1.12 Conclusion

The Logistic Regression model gives us an accuracy of 71% which means that among all of the positives that the model predicted, 71% were truly positive for the recall, out of the true positives and the false negatives, it was right 61% of the times. A way to improve the model would be to add more data on the missing values as It is oversampled on some classes, and less sampled on others

But, overall, we would train our model to improve on Its recall, as we should not have undetected diseases, as It might be fatal for the patient